

Table of Contents

Introduction.....	2
Accounts.....	3
Adaptive Echo.....	4
Block Cipher Modes of Operation.....	5
Communication Methods.....	6
Communication Model.....	7
Configuration Settings.....	8
Constant-Time Comparison Function.....	9
Echo Protocol.....	10
Echo Public Key Sharing.....	11
Electronic Mail.....	12
Electronic Mail Forward Secrecy.....	13
Encrypted and Authenticated Containers.....	14
File Encryption.....	15
Key Derivation.....	16
Hybrid System.....	17
Multiple Devices.....	18
Pass-through Devices.....	19
Poptastic.....	20
Public Key Infrastructure.....	21
Sessions.....	22
Socialist Millionaire Protocol.....	23
Two-Way Calling.....	24
Verifying Ownership of Public Keys.....	25
Wide Lanes.....	27
References.....	28

Introduction

Spot-On is an exploratory research project investigating a variety of communication and cryptographic algorithms. The software is composed of two separate applications, a multi-threaded kernel and a user interface. The two components are written in C++ and require the Qt framework as well as an assortment of libraries. Qt versions 4.8.x and Qt 5.x are supported. Qt version 4.7.x is loosely supported. The application is available on FreeBSD, Linux, OS X, OS/2, and Windows. The ARM architecture is also supported. Please note that the Echo algorithm and its name are not derived from Ernest J. H. Chang's 1982 Echo Algorithms: Depth Parallel Operations on General Graphs paper.

Accounts

Spot-On implements a plain, and perhaps original, two-pass mutual authentication protocol. The implementation is well-defined with or without SSL/TLS. The protocol is weakened if SSL/TLS is neglected, however. Please see the paragraph at the end of this section for additional details regarding the weakness. The Accounts procedure is as follows:

1. Binding endpoints are responsible for defining account information. During the account-creation process, an account may be designated for one-time use. Account names and account passwords each require at least 32 bytes of data.
2. After a network connection is established, a binding endpoint notifies the peer with an authentication request. The binding endpoint will terminate the connection if the peer has not identified itself within a fifteen-second window.
3. After receiving the authentication request, the peer responds to the binding endpoint. The peer submits the following information: $H_{\text{Hash Key}}(\text{Salt} \parallel \text{Time}) \parallel \text{Salt}$, where the Hash Key is a concatenation of the account name and the account password. The SHA-512 hash algorithm is presently used to generate the hash output. The Time variable has a resolution of minutes. The peer retains the salt value.
4. The binding endpoint receives the peer's information. Subsequently, it computes $H_{\text{Hash Key}}(\text{Salt} \parallel \text{Time})$ for all of the accounts that it possesses. If it does not discover an account, it increments Time by one minute and performs an additional search. If an account is discovered, the binding endpoint creates a message similar to the message created by the peer in the previous step and submits the information to the peer. The authenticated information is recorded. After a period of approximately 120 seconds, the information is destroyed.
5. The peer receives the binding endpoint's information and performs a similar validation process, including the analysis of the binding endpoint's salt. The two salt values must be distinct. The peer will terminate the connection if the binding endpoint has not identified itself within a fifteen-second window.

Please note that the Accounts system can be promoted by including an encryption key. The additional key will allow for finer time resolutions.

If SSL/TLS is not available, the protocol may be exploited. A relay station may record the values in the 3rd step and subsequently provide the information to the binding endpoint. The binding endpoint will therefore trust the foreign connection. The recording device may then seize the binding endpoint's response, the values in the 4th step, and provide the information to the peer. If the information is accurate, the peer will accept the binding endpoint's response.

Adaptive Echo

The Adaptive Echo is a complement to the Echo Protocol and substantiates the opinion that the Echo Protocol is a malleable method. Endpoints that bind multiple parties may optionally define Adaptive Echo tokens. Adaptive Echo tokens are composed of authentication and encryption keys as well as details about the choice algorithms. If configured, binding endpoints are able to permit or restrict information travel based on the content of the data. As an example, peers that are cognizant of a specific Adaptive Echo token will receive data from other cognizant peers whereas traditional peers will not. Binding endpoints therefore selectively-echo data.

The Adaptive Echo behaves as follows:

1. A binding endpoint defines an Adaptive Echo token. The information must be distributed securely.
2. A networked peer having the given Adaptive Echo token generates $H_{\text{Hash Key}}(E_{\text{Encryption Key}}(\text{Message} \parallel \text{Time})) \parallel E_{\text{Encryption Key}}(\text{Message} \parallel \text{Time})$ where the Encryption Key and Hash Key are derived from the Adaptive Echo token. The generated information is then submitted to the binding endpoint as Message \parallel Adaptive Echo Information.
3. The binding endpoint processes the received message to determine if the message is tagged with a known Adaptive Echo token. If the message is indeed tagged correctly, the Time value is inspected. If the Time value is within five seconds of the binding endpoint's local time, the message is considered correct and the peer's presence is recorded.
4. As the binding endpoint receives messages from other peers, it inspects the messages to determine if the messages have been tagged with Adaptive Echo tokens. This process creates a network of associated peers. Because peers themselves may be binding endpoints, the Adaptive Echo may be used to generate an artificial trust network.

Adaptive Echo is susceptible to eavesdropping. As an example, if a message that is tagged with an Adaptive Echo token should travel through one or more peers to reach a destination, the peers may record the message and subsequently replay the message to a binding peer. The replay must occur within the acceptance window of the message. Additionally, the binding endpoint's congestion control container must not already contain the message. If both conditions are met, the binding endpoint will consider the peer as trustworthy.

Block Cipher Modes of Operation

Spot-On uses CBC with CTS to provide confidentiality. The file encryption mechanism supports the GCM algorithm without the authenticity property that's provided by the algorithm. To provide authenticity, the application uses the encrypted-then-MAC (EtM) approach. The Encrypted and Authenticated Containers section provides more details.

Communication Methods

Spot-On supports Bluetooth, SCTP, TCP, and UDP (multicast and unicast) communication methods. For TCP-based communications, OpenSSL is supported. Spot-On distributes data with or without SSL/TLS. Please note that magnet distribution violates this principle and therefore requires SSL/TLS. Communications between the Spot-On Kernel and the Spot-On User Interface also require SSL/TLS via TCP. In essence, the application is generally transport-neutral.

Communication Model

Spot-On mostly assumes an asynchronous communication model. The Accounts, SMP, and Two-Way calling systems require responses.

Configuration Settings

Spot-On implements a defensive approach with respect to configuration settings. Shortly after the kernel and the user interface are started, important settings are reviewed and if necessary corrected. The potentially-adjusted values are stocked in global containers. Some methods also inspect critical values, adjusting them if necessary.

Constant-Time Comparison Function

Spot-On attempts to utilize constant-time byte comparison functions so as to avoid timing analysis.

Echo Protocol

Spot-On introduced the Echo Protocol. The Echo is a malleable protocol. That is, the protocol does not require rigid implementation details. Each model may adhere to their own peculiar obligations. The Echo functions on the elementary persuasion that information is dispersed over multiple or singular passages and channel endpoints evaluate the suitability of the received data. Because data may become intolerable, Spot-On implements its own congestion control algorithm. Received messages that meet some basic criteria are labeled and duplicates are discarded. Advanced models may define more sophisticated congestion-avoidance algorithms based upon their interpretations of the Echo.

Spot-On provides two modes of operation for the general Echo Protocol, Full Echo and Half Echo. The Full Echo permits absolute data flow. The Half Echo defines an agreement between two endpoints. Within this agreement, information from other endpoints is prohibited from traveling along the private channel.

Echo Public Key Sharing

The Echo Public Key Sharing construct was introduced by Mr. Fennel. It is an elegant compliment to the Echo Protocol. The concept may be summarized as follows:

1. A community is created. The community is defined by a pair of authentication and encryption keys. The keys are derived via the PBKDF2 function.
2. Public key pairs may be optionally exchanged via the community. Participants who subscribe to a well-defined community will automatically accept public key pairs from participants who have published their public keys to the respective community.

Electronic Mail

Spot-On provides two e-mail models for distributed e-mail. Endpoints may optionally define themselves as institutions or post offices, or both.

A brief description of e-mail institutions follows. E-mail institutions are artificially characterized by addresses and names. The information is not considered secret and several endpoints may identify themselves identically. It is the responsibility of an institution to accept subscribers, that is, public-key pairs. The data that an institution houses is stored in encrypted containers. Unlike physical institutions, Spot-On institutions are only allowed to read the signature portions of e-mail letters. The signatures allow verification of deposits and withdrawals.

The sole difference between e-mail institutions and e-mail post offices is that post offices require the distribution of public-key pairs.

Electronic Mail Forward Secrecy

This section briefly describes a two-step communication process for establishing forward secrecy in Spot-On e-mail. We assume a hybrid scheme with respect to public-key encryption.

Assumptions:

1. Permanent public key pairs have been exchanged correctly.
2. The respective kernels remain active during the exchange window.

Protocol:

1. Participant A generates an ephemeral public encryption key pair. The key pair's attributes are configurable. If the kernel is deactivated after the key pair is generated, the key pair is discarded and the protocol is terminated.
2. Participant A transmits the ephemeral public key to B. The key is encrypted with B's permanent encryption public key and optionally signed with A's permanent private signature key.
3. B receives the public key and optionally verifies the signature. If B requires a valid signature but one is not provided, the protocol is terminated.
4. B generates private authentication and encryption keys.
5. Using the ephemeral public encryption key, B transfers the keys to A. The complete bundle is encrypted with A's permanent encryption public key and optionally signed with B's permanent private signature key.
6. Participant A receives the bundle and optionally verifies the signature. If A requires a valid signature but one is not provided, the protocol is terminated.

The session keys generated in the fourth step may remain in use until one of the parties decides to establish new session keys.

Signatures are required over the Poptastic transport protocol.

Encrypted and Authenticated Containers

Some of the data that Spot-On retains locally is stored in encrypted and authenticated containers. CBC and CTS encryption modes are used with a variety of block ciphers. Encryption and authentication occur as follows:

1. If the size of the original data is less than the specified cipher's block size, the original data is re-sized such that its new size is identical to the cipher's block size. A zero-byte pad is applied.
2. Append the size of the original data to the original container.
3. Encrypt the augmented data via the selected cipher and specified mode.
4. Compute a keyed-hash of the encrypted container.
5. Concatenate the hash output with the encrypted data, $H_{\text{Hash Key}}(E_{\text{Encryption Key}}(\text{Data} \parallel \text{Size}(\text{Data}))) \parallel E_{\text{Encryption Key}}(\text{Data} \parallel \text{Size}(\text{Data}))$.

Spot-On also includes a mechanism for re-encoding data if new authentication and encryption keys are desired.

File Encryption

Spot-On includes a simple file encryption application. The PIN and secret are used for generating authentication and encryption keys. After the initial encryption key is derived and applied, subsequent encryption keys are derived via SHA-256. That is, $\text{Key}_{i+1} = \text{SHA-256}(\text{Key}_i)$. Several encryption and hash algorithms are supported as well as various encryption modes.

The format of the converted file is as follows:

Byte 0: 0 – unsigned file, 1 – signed file	Bytes 1 – 64: keyed hash or 64 bytes of zeros	Bytes 65 – EOF: encrypted data
--	---	--------------------------------

Key Derivation

Spot-On uses separate authentication and encryption keys for local data. The key-derivation process is as follows:

1. Generate a cryptographic salt. The size of the salt is configurable.
2. Derive a temporary key via the PBKDF2 function. The hash algorithm, iteration count, passphrase (question/answer), and salt are input parameters to the function. All of the aforementioned parameters are configurable.
3. Using the temporary key from the previous step, derive a new key via the PBKDF2 function. The previous parameters are also used, however, the temporary key replaces the passphrase (question/answer).
4. Separate the derived key into two distinct keys. The encryption key is N bytes long, where N is the recommended key size of the selected cipher. The remaining bytes compose the authentication key. The generated authentication key contains at least 512 bytes.

Hybrid System

Spot-On implements a hybrid system for authenticity and confidentiality. One portion of the system generates per-message authentication and encryption keys. These two keys are used for authenticating and encapsulating data. The two keys are encapsulated via the public-key portion of the system. The application also provides a mechanism for distributing session-like keys for data encapsulation. Again, the keys are encapsulated via the public-key system. An additional mechanism allows the distribution of session-like keys via previously-established private keys. Digital signatures are optionally applied to the data. As an example, please consider the following message: $E_{\text{Public Key}}(\text{Encryption Key} \parallel \text{Hash Key}) \parallel E_{\text{Encryption Key}}(\text{Data}) \parallel H_{\text{Hash Key}}(E_{\text{Public Key}}(\text{Encryption Key} \parallel \text{Hash Key}) \parallel E_{\text{Encryption Key}}(\text{Data}))$. The private-key authentication and encryption mechanism is identical to the procedure discussed in the Encrypted and Authenticated Containers section.

Multiple Devices

Spot-On allows identical representations of individual nodes across multiple devices. For example, one may configure a node on one device and copy the configuration to another device.

Pass-through Devices

Spot-On supports pass-through connections. Pass-through connections allow unaltered data travel. Please note that accounts are supported for such connections.

Poptastic

Version 0.17 of Spot-On introduced the Poptastic protocol. The new protocol allows participants to communicate via the IMAP, POP3, and SMTP protocols. The protocol provides a medium for near real-time conversations as well as traditional e-mail. Please note that the protocol does not support the two-way calling mechanism.

Public Key Infrastructure

Spot-On utilizes the libgcrypt and libntru libraries for permanent private and public key pairs. Presently, the application generates ten key pairs during the initialization process. Key generation is optional. Consequently, Spot-On does not require a public key infrastructure.

ElGamal, NTRU, and RSA encryption algorithms are supported. DSA, ECDSA, EdDSA, ElGamal, and RSA signature algorithms are supported. The OAEP and PSS schemes are used with RSA encryption and RSA signing, respectively.

Communications between nodes having diverse key types are well-defined if the nodes share common libgcrypt and libntru libraries.

Non-NTRU private keys are evaluated for correctness via the `gcry_pk_testkey()` function. Public keys must also meet some basic criteria such as including the public-key identifier.

Sessions

Excluding Forward Secrecy, SMP, accounts, and congestion control, Spot-On does not maintain a deliberate state system with respect to communications. Some portions of the system provide the user with state information, however, the information is not required for the health and stability of Spot-On communications.

Socialist Millionaire Protocol

As of version 0.19, Spot-On includes an asynchronous implementation of the Socialist Millionaire Protocol as defined by <https://otr.cypherpunks.ca/Protocol-v3-4.0.0.html>. Spot-On does not distribute zero-knowledge proofs during the various data exchanges. Also, Spot-On uses the SHA-512 of the secrets as the x and y components.

Assuming that Alice begins the exchange:

- Alice:
 1. Picks random exponents a_2 and a_3
 2. Sends Bob $g_{2a} = g_1^{a_2}$ and $g_{3a} = g_1^{a_3}$
- Bob:
 1. Picks random exponents b_2 and b_3
 2. Computes $g_{2b} = g_1^{b_2}$ and $g_{3b} = g_1^{b_3}$
 3. Computes $g_2 = g_{2a}^{b_2}$ and $g_3 = g_{3a}^{b_3}$
 4. Picks random exponent r
 5. Computes $P_b = g_3^r$ and $Q_b = g_1^r g_2^y$
 6. Sends Alice g_{2b} , g_{3b} , P_b and Q_b
- Alice:
 1. Computes $g_2 = g_{2b}^{a_2}$ and $g_3 = g_{3b}^{a_3}$
 2. Picks random exponent s
 3. Computes $P_a = g_3^s$ and $Q_a = g_1^s g_2^x$
 4. Computes $R_a = (Q_a / Q_b)^{a_3}$
 5. Sends Bob P_a , Q_a and R_a
- Bob:
 1. Computes $R_b = (Q_a / Q_b)^{b_3}$
 2. Computes $R_{ab} = R_a^{b_3}$
 3. Checks whether $R_{ab} == (P_a / P_b)$
 4. Sends Alice R_b
- Alice:
 1. Computes $R_{ab} = R_b^{a_3}$
 2. Checks whether $R_{ab} == (P_a / P_b)$
- If everything is done correctly, then R_{ab} should hold the value of (P_a / P_b) times $(g_2^{a_2 b_3})^{(x-y)}$, which means that the test at the end of the protocol will only succeed if $x == y$. Further, since $g_2^{a_2 b_3}$ is a random number not known to any party, if x is not equal to y , no other information is revealed.

Two-Way Calling

Spot-On implements a plain two-pass key-distribution system. The protocol is defined as follows:

1. A peer generates 128-bit AES and 256-bit SHA-512 keys via the system's cryptographic random number generator.
2. Using the destination's public key, the peer encapsulates the two keys via the hybrid cryptographic system.
3. The destination peer receives the data, records it, and generates separate keys as in step 1.
4. The destination peer transmits the encapsulated keys to the originating peer as in step 2.

Once the protocol is executed, the two peers shall possess identical authentication and encryption keys. Please note that duplicate half-keys are allowed.

Verifying Ownership of Public Keys

A shared public key bundle includes the encryption public key, a signature of the encryption public key, the signature public key, and a signature of the signature public key. Ownership of NTRU encryption keys are not verified. Signatures are generated as follows:

1. The encryption or signature public key is gathered and stored in *data_t*.
2. The SHA-512 digest of the public key is computed and stored in *hash*.
3. For DSA, Elliptic Curve DSA, and ElGamal public keys:
 1. `gcry_sexp_build(&data_t, 0,`
 `"(data (flags raw)(value %m))",`
 `hash_t)` is computed.
4. For Edward Curve DSA public keys:
 1. `gcry_sexp_build(&data_t, 0,`
 `"(data (flags eddsa)(hash-algo sha512)"`
 `"(value %b))",`
 `hash.length(),`
 `hash.constData())` is computed.
5. For RSA public keys:
 1. `gcry_sexp_build(&data_t, 0,`
 `"(data (flags pss)(hash sha512 %b)"`
 `"(random-override %b))",`
 `hash.length(),`
 `hash.constData(),`
 `random.length(),`
 `random.constData())` is computed, where the variable *random* contains 20 bytes of random data.
6. Finally, `gcry_pk_sign(&signature_t, data_t, key_t)` is computed. The contents of *signature_t* are extracted via `gcry_sexp_sprint()` and stored in a byte array. The variable *key_t* contains either the public encryption key or the public signature key.

Signatures are verified as follows:

1. The SHA-512 digest of the provided public key is computed and stored in *hash*.
2. For DSA, Elliptic Curve DSA, and ElGamal public keys:
 1. `gcry_sexp_build(&data_t, 0,`
 `"(data (flags raw)(value %m))",`

`hash_t`) is computed. The contents of *hash_t* are populated via `gcry_mpi_scan()`.

3. For Edward Curve DSA public keys:

1. `gcry_sexp_build(&data_t, 0,`
 `"(data (flags eddsa)(hash-algo sha512)"`
 `"(value %b))",`
 `hash.length(),`
 `hash.constData())` is computed.

4. For RSA public keys:

1. `gcry_sexp_build(&data_t, 0,`
 `"(data (flags pss)(hash sha512 %b)"`
 `"(random-override %b))",`
 `hash.length(),`
 `hash.constData(),`
 `random.length(),`
 `random.constData())` is computed, where the variable *random* contains 20 bytes of random data.

5. Finally, `gcry_pk_verify(signature_t, data_t, key_t)` is computed. If `gcry_pk_verify()` returns zero, the signature is considered valid. The contents of *signature_t* are populated via `gcry_sexp_new()` using the provided signature array.

Wide Lanes

One of the many obligations of a Spot-On-Kernel process is to receive, process, and forward data to one or more nodes. The mechanism that performs this task is similar to both a network hub and a network switch. Wide Lanes allow node operators to assign listener lane widths. Let's consider a basic example, a listener having a lane width of 20,000 bytes. The kernel, if necessary, will forward packets via the listener's clients if the sizes of the forwarded packets do not exceed 20,000 bytes. Optionally, clients may negotiate different lane widths with their peers.

References

http://world.std.com/~dtd/sign_encrypt/sign_encrypt7.html
https://en.wikipedia.org/wiki/Authenticated_encryption
https://en.wikipedia.org/wiki/Block_cipher_mode_of_operation
https://en.wikipedia.org/wiki/Ciphertext_stealing
https://en.wikipedia.org/wiki/Hybrid_cryptosystem
<https://en.wikipedia.org/wiki/NTRU>
https://en.wikipedia.org/wiki/Optimal_asymmetric_encryption_padding
https://en.wikipedia.org/wiki/PKCS_1
https://en.wikipedia.org/wiki/Padding_%28cryptography%29
https://en.wikipedia.org/wiki/Stream_Control_Transmission_Protocol
https://en.wikipedia.org/wiki/Timing_attack
<https://otr.cypherpunks.ca/Protocol-v3-4.0.0.html>
<https://www.cs.jhu.edu/~astubble/dss/ae.pdf>